

Towards Efficient String Processing of Annotated Events

David Woods
ADAPT Centre
Trinity College Dublin, Ireland
dwoods@tcd.ie

Tim Fernando, Carl Vogel
Computational Linguistics Group
Trinity Centre for Computing and Language Studies
School of Computer Science and Statistics
Trinity College Dublin, Ireland
tim.fernando@tcd.ie, vogel@tcd.ie

Abstract

This paper explores the use of strings as models to effectively represent event data such as might be found in a document annotated with ISO-TimeML. We describe the translation of such data to strings, as well as a number of operations, such as superposition, which may be used to manipulate these strings in order to infer new information. Some advantages and limitations of the operations are discussed, along with how the latter might be controlled by imposing constraints on the models, in order to avoid over-generation when superposing the string models. In particular, we look at how Allen Relations, which might be extracted from a document annotated with ISO-TimeML, can be translated simply to strings, and create a useful set of constraints in doing so.

1 Introduction

This paper explores the use of strings as models to effectively represent event data such as might be found in a document annotated with ISO-TimeML. It is described how such data may be simply translated to strings, and how to infer information through operations on these strings. Strings are basic computational entities that can be more readily manipulated by machines than the infinite models of predicate logic. Finite sets of strings serve as finite models.

Given a finite set A of fluents (predicates with an associated temporality), and treating sets as symbols so that we might construct a representational string from them (as in Fernando, 2016), a string $s = \alpha_1 \cdots \alpha_n$ of subsets α_i of A can be construed as a finite model consisting of n moments of time $i \in \{1, \dots, n\}$ with α_i specifying all fluents (in A) that (as unary predicates) hold simultaneously at i .

Throughout this paper, a fluent $a \in A$ will be understood as naming an event, and the powerset 2^A of A will serve as an alphabet $\Sigma = 2^A$ of an *event-string* $s \in \Sigma^+$. Such strings are finite models of Monadic Second Order logic, and are amenable to finite state methods. We will further restrict them in Section 6, with a focus on using Allen Relations adopted in ISO-TimeML, in order to analyse inference over a finite search space.

An event-string $\alpha_1 \cdots \alpha_n$ is read from left to right chronologically, so that any predicates which hold at the moment at index i are understood to have held before another moment indexed by j if and only if $i < j$. The precise duration of each moment is taken as unimportant in the current discussion, and thus the strings model an inertial world, whereby *change* is the only mark of progression from one moment to the next – “But neither does time exist without change” (Aristotle, *Physics IV*). Thus, if $\alpha_i = \alpha_{i+1}$ for any $1 \leq i < n$, then either α_i or α_{i+1} may be safely deleted from s without affecting the interpretation of the string, as the remaining symbol is simply taken as representing a longer moment. This operation of removing repetition from the event-string is known as *block compression* (Fernando, 2016). The inverse of this process introduces repeated elements in an event-string for greater flexibility in manipulating strings. These operations are detailed in Section 3.

We see that strings may provide useful finite models for event data, once sufficiently constrained. This is in order to avoid a large combinatorial blow-up when reconciling information from different strings.

In the following section, we see an example from ISO-TimeML which we might convert to an event-string, and why it would be useful to do so. Sections 3 and 4 provide the formal string operations used to manipulate and combine information from event-strings. How these operations are useful in regards to Allen Relations is described in Section 5, and Section 6 explores constraints to prevent malformed event-strings. Finally, Section 7 shows how the framework can be applied to data from an ISO-TimeML document.

2 Motivation in ISO-TimeML

ISO-TimeML (Pustejovsky et al., 2010) is a standard markup language used for the annotation of events (and their interrelations) in texts. Of particular interest to us here are the TLINK elements, which indicate the relations between pairs of fluents found in the document. Though not every fluent will necessarily be linked with another in this manner, a majority will be. The TIMEBANK Corpus (Pustejovsky et al., 2003) provides a large number of documents annotated using the 1.1 TimeML standard (a predecessor to ISO-TimeML), which we may extract TLINKs from.

According to the markup specification (TimeML Working Group, 2005), a TLINK is required to have the following attributes: either a `timeID` or `eventInstanceId` attribute, referring to some fluent in the text, as well as either a `relatedToTime` or `relatedToEventInstance` attribute, which will refer to another fluent, and also a `relType` attribute, declaring the relation between the two fluents. Other attributes are optional and not relevant to the current discussion.

In order to give a more concrete understanding, let us take a small fragment from an ISO-TimeML document, which will give us three TLINK nodes:

- (1a) `<TLINK lid="l2" relType="IS_INCLUDED" eventInstanceId="ei1" relatedToTime="t1" origin="USER"/>`
- (1b) `<TLINK lid="l9" relType="IS_INCLUDED" timeID="t1" relatedToEventInstance="ei9" origin="USER"/>`
- (1c) `<TLINK lid="l10" relType="BEFORE" eventInstanceId="ei9" relatedToEventInstance="ei10" origin="USER"/>`

The value of `relType` will correspond to exactly one of the relations described in Allen (1983), though it should be noted that some of these Allen Relations can correspond to multiple `relTypes` (for example, `IDENTITY` and `SIMULTANEOUS` are both covered by the Allen Relation *equal*). The other attributes (`lid`, `origin`) in each TLINK may be ignored for now.

We can represent the information in (1a)–(1c) in predicate logic as follows:

- (2a) $Includes(t1, ei1)$
- (2b) $Includes(ei9, t1)$
- (2c) $Before(ei9, ei10)$

where *Includes* (the inverse of `IS_INCLUDED`) and *Before* are binary relations, corresponding to the Allen Relations *during* (inverse) and *before*, respectively (see Section 5).

An issue with this representation is that the full picture of the chronological sequence of events is not intuitively obvious from the three predicates. It is possible to create a set of inference rules (see Setzer et al., 2005) which allow for drawing new conclusions from the given information – for example, that the predicate *Before*($t1, ei10$) also holds. However, the amount of information that can be obtained from any single binary relation is relatively small, and reasoning about more than two fluents at once can require the conjunction of multiple relations.

Using an event-string, we may include all of the above information in a single, readable string, which allows us to reason about the relations between any number of fluents. Additionally, we provide facilities to infer new relations from this event-string (see Section 6).

3 Superposition and Block Compression

In order to usefully collect information from multiple strings into a single string, we define here the operation of *superposition*. With two strings s and s' of the same length n built from an alphabet Σ , the powerset of some fixed set A , the superposition $s \& s'$ of s and s' is their componentwise union:

$$(3) \quad \alpha_1 \cdots \alpha_n \& \alpha'_1 \cdots \alpha'_n := (\alpha_1 \cup \alpha'_1) \cdots (\alpha_n \cup \alpha'_n)$$

For convenience of notation, we will use boxes rather than curly braces $\{ \}$ to represent sets in Σ , such that each symbol α in a string s corresponds to exactly one box. For example, with $a, b, c, d \in A$:

$$(4) \quad \boxed{a|c} \& \boxed{b|d} = \boxed{a,b|c,d} \in \Sigma^2$$

Extending superposition to languages L and L' over the same alphabet is a simple matter of collecting the superpositions of strings of equal length from each language:

$$(5) \quad L \& L' := \bigcup_{n \geq 0} \{s \& s' \mid s \in L \cap \Sigma^n \text{ and } s' \in L' \cap \Sigma^n\}$$

For example, $L \& \square^* = L$. If L and L' are regular languages computed by finite automata with transitions \rightarrow and \rightarrow' , then the superposition $L \& L'$ is a regular language computed by a finite automaton with transitions \Rightarrow formed by running \rightarrow and \rightarrow' in lockstep according to the rule

$$(6) \quad \frac{q \xrightarrow{\alpha} r \quad q' \xrightarrow{\alpha'} r'}{(q, q') \xRightarrow{\alpha \cup \alpha'} (r, r')}$$

A disadvantage of this operation is that it requires the string operands to be of equal length, which is an overly specific case. In order to generalise this procedure to strings of arbitrary lengths, we may manipulate the strings to move away from the synchrony of the lockstep procedure. One such manipulation is that we can cause a string $s = \alpha_1 \cdots \alpha_n$ to *stutter* such that $\alpha_i = \alpha_i + 1$ for some integer $0 < i < n$. For example, $\boxed{a|a|a|c|c}$ is a stuttering version of $\boxed{a|c}$. If a string does not stutter, it is *stutterless*, and we can transform a stuttering string to this state by using “block compression”:

$$(7) \quad \text{bc}(s) := \begin{cases} s & \text{if } \text{length}(s) \leq 1 \\ \text{bc}(\alpha s') & \text{if } s = \alpha \alpha s' \\ \alpha \text{bc}(\alpha' s') & \text{if } s = \alpha \alpha' s' \text{ with } \alpha \neq \alpha' \end{cases}$$

This function can be applied multiple times to a string, but the output will not change after the first application: $\text{bc}(\text{bc}(s)) = \text{bc}(s)$. We can also use the inverse of this function to generate infinitely many stuttering strings:

$$(8) \quad \text{bc}^{-1}(\boxed{a|c}) = \{\boxed{a|c}, \boxed{a|a|c}, \boxed{a|c|c}, \boxed{a|a|c|c}, \dots\}$$

We can say that any of the strings generated by this inverse block compression are *bc-equivalent*. Precisely, a string s' is bc-equivalent to a string s iff $s' \in \text{bc}^{-1}\text{bc}(s)$.

We can now define the *asynchronous superposition* $s \&_* s'$ of strings s and s' as the (provably) *finite* set obtained by block compressing the *infinite* language generated by superposing the strings which are bc-equivalent to s and s' :

$$(9) \quad s \&_* s' := \{\text{bc}(s'') \mid s'' \in \text{bc}^{-1}\text{bc}(s) \& \text{bc}^{-1}\text{bc}(s')\}$$

For example, $\boxed{a|c} \&_* \boxed{b|d}$ will comprise three strings:

$$(10) \quad \{\boxed{a,b|c,d}, \boxed{a,b|a,d|c,d}, \boxed{a,b|b,c|c,d}\}$$

In order to avoid generating all possible strings when using the inverse block compression, we introduce an upper bound to the length of the strings which will be superposed. It can be shown that with two strings of length n and n' , the longest bc-unique string (one which has no shorter bc-equivalent strings) produced through asynchronous superposition will be of length $n + n' - 1$.

4 Upper Bound on Asynchronous Superposition

The aforementioned bound on the length of superposed strings can be established as follows. For all $s, s' \in \Sigma^*$, we define a finite set $s \hat{\&} s'$ of strings over Σ with enough of the strings in $\text{bc}^{-1}\text{bc}(s) \& \text{bc}^{-1}\text{bc}(s')$ to form $s \&_* s'$. The definition proceeds by induction on s and s' , with

$$(11a) \quad \epsilon \hat{\&} \epsilon := \{\epsilon\}$$

$$(11b) \quad \epsilon \hat{\&} s := \emptyset \quad \text{for } s \neq \epsilon$$

$$(11c) \quad s \hat{\&} \epsilon := \emptyset \quad \text{for } s \neq \epsilon$$

and for all $\alpha, \alpha' \in \Sigma$,

$$(12) \quad \alpha s \hat{\&} \alpha' s' := \{(\alpha \cup \alpha')s'' \mid s'' \in (\alpha s \hat{\&} s') \cup (s \hat{\&} \alpha' s') \cup (s \hat{\&} s')\}$$

Note that a string in $s \hat{\&} s'$ might stutter, even if neither of the operands s or s' do (e.g. $\boxed{a, c} \boxed{a, c} \in \boxed{a} \boxed{c} \hat{\&} \boxed{c} \boxed{a}$). However, it can be made stutterless through block compression.

Proposition 1. *For all $s, s' \in \Sigma^+$ and all $s'' \in s \hat{\&} s'$,*

$$(13) \quad \text{length}(s'') \leq \text{length}(s) + \text{length}(s') - 1$$

Proposition 2. *For all $s, s' \in \Sigma^+$,*

$$(14) \quad s \hat{\&} s' \subset \text{bc}^{-1}\text{bc}(s) \& \text{bc}^{-1}\text{bc}(s')$$

and

$$(15) \quad \{\text{bc}(s'') \mid s'' \in s \hat{\&} s'\} = s \&_* s'$$

Now, for any integer $k > 0$ and string $s = \alpha_1 \cdots \alpha_n$ over Σ , we introduce a new function pad_k which will generate the set of strings with length k which are bc -equivalent to s :

$$(16a) \quad \text{pad}_k(\alpha_1 \cdots \alpha_n) := \alpha_1^+ \cdots \alpha_n^+ \cap \Sigma^k$$

$$(16b) \quad = \{\alpha_1^{k_1} \cdots \alpha_n^{k_n} \mid k_1, \dots, k_n \geq 1 \text{ and } \sum_{i=1}^n k_i = k\}$$

$$(16c) \quad \subset \text{bc}^{-1}\text{bc}(\alpha_1 \cdots \alpha_n)$$

For example, $\text{pad}_4(\boxed{a} \boxed{c})$ will generate $\{\boxed{a} \boxed{a} \boxed{a} \boxed{c}, \boxed{a} \boxed{a} \boxed{c} \boxed{c}, \boxed{a} \boxed{c} \boxed{c} \boxed{c}\}$. We can use this new function in our calculation of asynchronous superposition, to limit the generation of strings from the inverse block compression step. Since we know from Proposition 1 that the maximum possible length we might need is $n + n' - 1$, we can use this value in the pad function to just generate the strings of that length, giving us a new definition of asynchronous superposition:

Corollary 3. *For any $s, s' \in \Sigma^+$ with nonzero lengths n and n' respectively,*

$$(17) \quad s \&_* s' = \{\text{bc}(s'') \mid s'' \in \text{pad}_{n+n'-1}(s) \& \text{pad}_{n+n'-1}(s')\}$$

Neither $s \hat{\&} s'$ nor $\text{pad}_{n+n'-1}(s) \& \text{pad}_{n+n'-1}(s')$ need be a subset of the other, even though, under the assumptions of Corollary 3, both sets block compress to $s \&_* s'$.

5 Event Representation

Now we may use asynchronous superposition to generate the 13 strings in $\boxed{e} \&_* \boxed{e'}$, each of which corresponds to one of the unique interval relations in Allen (1983). No more than one of these relations may hold between any two fluents, and thus each of the 13 generated event-strings exists in a distinct possible “world”. We use the empty box $\boxed{}$ as a string of length 1 (not to be confused with the empty string ϵ , which is length 0) to bound events, allowing us to represent the fact that they are finite – they have a beginning and ending point. It is prudent to assume that we will deal only with finite event data, such that there are no fluents which do not have both an associated start-point and end-point. If such a non-finite fluent without a beginning and ending were to occur, it could trivially appear in every position in the event-string.

The bounding boxes represent the time before and after the event occurs, during which no other fluents $a \in A$ are mentioned. The event-strings associated with the Allen Relations are laid out below:

$e = e'$	$\boxed{e, e'}$	equal
$e \text{ s } e'$	$\boxed{e, e'} \boxed{e'}$	starts
$e \text{ si } e'$	$\boxed{e, e'} \boxed{e}$	starts (inverse)
$e \text{ f } e'$	$\boxed{e'} \boxed{e, e'}$	finishes
$e \text{ fi } e'$	$\boxed{e} \boxed{e, e'}$	finishes (inverse)
$e \text{ d } e'$	$\boxed{e'} \boxed{e, e'} \boxed{e'}$	during
$e \text{ di } e'$	$\boxed{e} \boxed{e, e'} \boxed{e}$	during (inverse)
$e \text{ o } e'$	$\boxed{e} \boxed{e, e'} \boxed{e'}$	overlaps
$e \text{ oi } e'$	$\boxed{e'} \boxed{e, e'} \boxed{e}$	overlaps (inverse)
$e \text{ m } e'$	$\boxed{e} \boxed{e'}$	meets
$e \text{ mi } e'$	$\boxed{e'} \boxed{e}$	meets (inverse)
$e < e'$	$\boxed{e} \boxed{e'}$	before
$e > e'$	$\boxed{e'} \boxed{e}$	after

These Allen Relations are included in the attributes of ISO-TimeML, as types of relation annotated by TLINKs (though some relations are named slightly differently). By extracting the TLINKs from an annotated document, and translating them to our event-string representation (see Section 7), we may begin to reason about the relationships between annotated events which do not have an associated TLINK in the markup. For example, the document may give us a relation between events e and e' , and another relation between e' and e'' , and from this we may infer the possible relations between e and e'' .

As asynchronous superposition is commutative and associative, we may superpose arbitrary numbers of event-strings: $s_1 \&_* \dots \&_* s_n$. We can show that superposing n unconstrained bounded event-strings will generate strings of maximum length $2n + 1$.¹ Note, however, that superposing even a relatively small number of unconstrained bounded events leads to a massive combinatorial blow-up in the number of

¹The proof is by induction:

Let each string to be superposed $s_i \in \{s_1, \dots, s_n\}$ be $\boxed{e_i}$, with each $e_i \in A$.

For $n = 2$: $s_1 \&_* s_2$.

From Proposition 1, the maximum length of the result is $3 + 3 - 1 = 5 = 2(2) + 1$.

We assume true for $n = p$, thus the maximum length of $s_1 \&_* \dots \&_* s_p$ is $2(p) + 1$.

Next, we prove for $n = p + 1$: $s_1 \&_* \dots \&_* s_{p+1} = s_1 \&_* \dots \&_* s_p \&_* s_{p+1} = s_{1\dots p} \&_* s_{p+1}$.

From Proposition 1, the maximum length of the result is $(2(p) + 1) + 3 - 1 = 2(p + 1) + 1$.

Thus true for $p + 1$, and by induction, true for any $n \geq 2$.

outcomes, or possible worlds, as each event-string generated from one superposition (e.g. $s_1 \&_* s_2$) will in turn be superposed with each generated from another (e.g. $s_3 \&_* s_4$). Additionally, with each fluent, the maximum possible length of the strings grows, meaning a larger set of strings will be generated at the *pad* stage. The combinatorics for n unconstrained bounded events are show in Table 1, up to $n = 5$. Clearly, simply superposing bounded events in this manner is not feasible, given the huge growth in the

Number of unconstrained bounded events	Number of outcomes from asynchronous superposition
2	13
3	409
4	23917
5	2244361

Table 1: Outcomes of Asynchronous Superposition with n Inputs

number of possible outcomes, and it is unreasonable to expect that any given document should contain five or fewer events. In order to avoid generating such a large number of computed event-strings, it is necessary to add constraints where appropriate to limit the strings that may be considered allowable for a particular context.

Interestingly, because each unconstrained bounded event-string \boxed{e} contains exactly one fluent, we may determine the maximum possible length of a string generated by superposition, $2n + 1$, from the size of the set A of fluents, where $n = |A|$. By keeping track of $|A|$, we ensure that the length of the string will always be finite, opening up the possibility of using methods from constraint satisfaction, exploiting the finite search space.

6 Constraints on Event-Strings

Two approaches to constraints may be implemented, which are not mutually exclusive. The first is to prevent unwanted strings from being generated, based on the nature of the operand strings, and the second is to remove disallowed strings from the set of outputs. The former approach is preferred from a computational standpoint, as there is less data to store and process. For either, we define some properties of what we may consider to be a *well-formed event-string*.

We assume that every fluent we encounter has exactly one beginning and one ending – that is, that events do not *resume* once they have ended. Events of the same type may stop and start frequently, but by assuming that every instance of an event will have a uniquely identifying fluent, we can discard any strings which feature such a resumption.² In this way, fluents are *interval-like*. We define the function ρ_X on strings of sets to component-wise intersect with X for any $X \subseteq A$ (Fernando, 2016):

$$(18) \quad \rho_X(\alpha_1 \cdots \alpha_n) := (\alpha_1 \cap X) \cdots (\alpha_n \cap X)$$

Applying block compression to an event-string which has been reduced with $\rho_{\{a\}}$ should produce a single string: \boxed{a} . For example, with $a, b \in A$:

$$(19) \quad \text{lc}(\rho_{\{a\}}(\boxed{a} \mid a, b \mid b)) = \boxed{a}$$

Additionally, fluents may be referred to multiple times by different TLINKs in an annotated document, and we assume that they will be *consistent* within the context of that document *i.e.* if a relation holds between e and e' , and a relation holds between e' and e'' , then both instances of e' refer to the same fluent. In this case, if a relation also holds between e and e'' , then this relation should not contradict the

²We adopt simplifying assumptions made in Allen Relations, though it should be noted that the distinction between event instances and event types (see Fernando, 2015) is not imposed by the event-string framework itself, allowing for discontinuous events (such as *judder*) in future work.

other two relations. For example, if $e > e'$ and $e' > e''$, then it should be impossible for a well-formed event-string to also have the relation $e < e''$, as this would break the interval-like fluent constraint mentioned above.

These last two points are interesting in particular, as they lead to a specific kind of superposition between strings $s, s' \in \Sigma^+$ when some symbol $\alpha \in s$ is equal to some other symbol $\alpha' \in s'$. In this scenario, the symbols must unify when superposing the strings, in order to create a well-formed event-string in accordance with the above two constraints. To achieve this, when a symbol α in s is also present in s' , and the asynchronous superposition of these strings is desired, padding is carried out as normal, but superposition is only permitted of those results of padding in which the indices of the matching symbols are equal. To do otherwise would permit event-strings which are not well-formed.

Allen (1983) gives a transitivity table showing the inferred possible relations between two events a and c , given the relation between each and an intermediary event, b . Each cell of the table shows simply the symbol which represent the binary relation – we may improve on the readability of this by showing explicitly the well-formed event-string(s) formed by the asynchronous superposition in each case. A fragment of the entire table is shown in Table 2 below:

	“before” <table><tr><td>b</td><td>c</td></tr></table>	b	c	“during” <table><tr><td>c</td><td>b, c</td><td>c</td></tr></table>	c	b, c	c	“meets” <table><tr><td>b</td><td>c</td></tr></table>	b	c	“starts” <table><tr><td>b, c</td><td>c</td></tr></table>	b, c	c																								
b	c																																				
c	b, c	c																																			
b	c																																				
b, c	c																																				
“before” <table><tr><td>a</td><td>b</td></tr></table>	a	b	<table><tr><td>a</td><td>b</td><td>c</td></tr></table>	a	b	c	<table><tr><td>a</td><td>c</td><td>b, c</td><td>c</td></tr><tr><td>a</td><td>a, c</td><td>c</td><td>b, c</td><td>c</td></tr><tr><td>a</td><td>c</td><td>b, c</td><td>c</td></tr><tr><td>c</td><td>a, c</td><td>c</td><td>b, c</td><td>c</td></tr><tr><td>a, c</td><td>c</td><td>b, c</td><td>c</td></tr></table>	a	c	b, c	c	a	a, c	c	b, c	c	a	c	b, c	c	c	a, c	c	b, c	c	a, c	c	b, c	c	<table><tr><td>a</td><td>b</td><td>c</td></tr></table>	a	b	c	<table><tr><td>a</td><td>b, c</td><td>c</td></tr></table>	a	b, c	c
a	b																																				
a	b	c																																			
a	c	b, c	c																																		
a	a, c	c	b, c	c																																	
a	c	b, c	c																																		
c	a, c	c	b, c	c																																	
a, c	c	b, c	c																																		
a	b	c																																			
a	b, c	c																																			
“during” <table><tr><td>b</td><td>a, b</td><td>b</td><td>c</td></tr></table>	b	a, b	b	c	<table><tr><td>b</td><td>a, b</td><td>b</td><td>c</td></tr></table>	b	a, b	b	c	<table><tr><td>c</td><td>b, c</td><td>a, b, c</td><td>b, c</td><td>c</td></tr></table>	c	b, c	a, b, c	b, c	c	<table><tr><td>b</td><td>a, b</td><td>b</td><td>c</td></tr></table>	b	a, b	b	c	<table><tr><td>b, c</td><td>a, b, c</td><td>b, c</td><td>c</td></tr></table>	b, c	a, b, c	b, c	c												
b	a, b	b	c																																		
b	a, b	b	c																																		
c	b, c	a, b, c	b, c	c																																	
b	a, b	b	c																																		
b, c	a, b, c	b, c	c																																		
“meets” <table><tr><td>a</td><td>b</td></tr></table>	a	b	<table><tr><td>a</td><td>b</td><td>c</td></tr></table>	a	b	c	<table><tr><td>a</td><td>a, c</td><td>b, c</td><td>c</td></tr><tr><td>c</td><td>a, c</td><td>b, c</td><td>c</td></tr><tr><td>a, c</td><td>b, c</td><td>c</td></tr></table>	a	a, c	b, c	c	c	a, c	b, c	c	a, c	b, c	c	<table><tr><td>a</td><td>b</td><td>c</td></tr></table>	a	b	c	<table><tr><td>a</td><td>b, c</td><td>c</td></tr></table>	a	b, c	c											
a	b																																				
a	b	c																																			
a	a, c	b, c	c																																		
c	a, c	b, c	c																																		
a, c	b, c	c																																			
a	b	c																																			
a	b, c	c																																			
“starts” <table><tr><td>a, b</td><td>b</td></tr></table>	a, b	b	<table><tr><td>a, b</td><td>b</td><td>c</td></tr></table>	a, b	b	c	<table><tr><td>c</td><td>a, b, c</td><td>b, c</td><td>c</td></tr></table>	c	a, b, c	b, c	c	<table><tr><td>a, b</td><td>b</td><td>c</td></tr></table>	a, b	b	c	<table><tr><td>a, b, c</td><td>b, c</td><td>c</td></tr></table>	a, b, c	b, c	c																		
a, b	b																																				
a, b	b	c																																			
c	a, b, c	b, c	c																																		
a, b	b	c																																			
a, b, c	b, c	c																																			

Table 2: Fragment of Allen Transitivity Table using event-strings

Here and in the original table, only three events are mentioned: a , b , and c . We can see that the asynchronous superposition of an event-string $s_{a,b}$ mentioning a and b with an event-string $s_{b,c}$ mentioning b and c gives a language L of event-strings mentioning all three events. Applying the reduct $\rho_{\{a,b\}}$ to any string in L (and block compressing the result) should give back exactly $s_{a,b}$, and likewise applying $\rho_{\{b,c\}}$ to any string in L should give back exactly $s_{b,c}$. It should, in theory, be possible to generalise this to any number of events, ensuring the same level of readability by using event-strings. Then for an event-string of arbitrary length featuring any number of fluents, we can apply the reduct $\rho_{\{e,e'\}}$ and block compress the result to obtain the Allen Relation between fluents e and e' by comparing with the event-strings corresponding to each Allen Relation, laid out in Section 5.

Finally, we may also introduce further constraints if external information is available, and these might be simply intersected with the result of a superposition: $(s \&_* s') \cap C$, where C represents the constraints to be applied, for example, “ e is among the first events to occur in the string s ” (true iff $s \& \boxed{e}^* = s$). This allows for extension beyond Allen Relations in the future.

7 Application to TLINKs

As mentioned in Section 5, attempting to generate all of the possible worlds becomes difficult when using just the unconstrained bounded event-strings alone, as there are just too many (rarely, if ever, featuring five fluents or fewer). Instead, we begin by looking at just those fluents which are linked to another by Allen Relation, which we may extract from the TLINKs in a document annotated with ISO-TimeML, as noted in Section 2.

As each relation corresponds exactly to one possible model, we translate the TLINKs immediately to the appropriate event-strings, and superpose these according to the constraints mentioned in Section 6. This allows us to avoid simply superposing based on the fluents, and bypasses having to generate the initial 13 possibilities. In this way, we may generate a much smaller set of possible outcomes from a larger number of bounded events.

We can now rewrite the information given in the three TLINKs (1a)–(1c) as the following event-strings:³

$$\begin{aligned}
 (20a) \quad & \boxed{t1} \boxed{ei1, t1} \boxed{t1} \\
 (20b) \quad & \boxed{ei9} \boxed{t1, ei9} \boxed{ei9} \\
 (20c) \quad & \boxed{ei9} \boxed{ei10}
 \end{aligned}$$

We may asynchronously superpose these event-strings while respecting the established constraints in order to generate a new event-string which contains all of the information from each of the inputs:

$$(21) \quad \boxed{ei9} \boxed{t1, ei9} \boxed{ei1, t1, ei9} \boxed{t1, ei9} \boxed{ei9} \boxed{ei10}$$

A clear advantage here is the compact, readable representation. Furthermore, we can infer the Allen Relation between any two fluents in this event-string s by applying a reduct and block compressing the result. For example, to infer the relationship between $ei1$ and $ei10$, we obtain $s' = \text{bc}(\rho_{\{ei1, ei10\}}(s)) = \boxed{ei1} \boxed{ei10}$, from which we can conclude that the Allen Relation $ei1$ *before* $ei10$ holds.

A drawback here is that for this to be effective, it relies on the events being heavily constrained by their interrelations. If there are too few TLINKs relative to the number of events, we still run into the problem of combinatorial explosion. What's more, as seen in Table 2, some asynchronous superpositions will still generate multiple disjoint possibilities (such as a *before* b with b *during* c), which will also impact the combinatorial problem.

An additional issue in computation of the superposition of events arises as multiple superposition operations are carried out in sequence, meaning unordered data may lead to a much less efficient calculation of final results. For example, $\boxed{a} \boxed{b} \&_* \boxed{b} \boxed{c} \&_* \boxed{c} \boxed{d}$ and $\boxed{a} \boxed{b} \&_* \boxed{c} \boxed{d} \&_* \boxed{b} \boxed{c}$ should produce the same, single output: $\boxed{a} \boxed{b} \boxed{c} \boxed{d}$, which they do. However, due to the respective orderings, the first sequence will arrive at that conclusion much faster as $\boxed{a} \boxed{b} \&_* \boxed{b} \boxed{c}$ has one possible outcome ($\boxed{a} \boxed{b} \boxed{c}$), which can immediately be asynchronously superposed with $\boxed{c} \boxed{d}$ to produce the final output. However, $\boxed{a} \boxed{b} \&_* \boxed{c} \boxed{d}$ has 321 possible outcomes, each of which must be individually asynchronously superposed with $\boxed{b} \boxed{c}$, only to come to the same conclusion, as only one of these results is well-formed.

One potential way to work around this pitfall is a grouping and ordering stage, where initially only events linked by some relation may be superposed, and only after the operand strings have been sorted to some optimal order, whereby the event-strings with the most shared fluents are grouped. It may be prudent to only perform superposition at all on event-strings which may be linked through one or more

³Using <https://www.scss.tcd.ie/~dwoods/timeml/> to quickly extract TLINKs from an ISO-TimeML document and translate them to event-strings. A non-trivial extension of this program which computes results of superposition is in the works.

relations or shared fluents. In this way, new, underspecified events may be formed from the output strings. Consider the scenario with $e_1, \dots, e_8 \in A$, and the following Allen Relations:

- (22a) $e_1 < e_4$
- (22b) $e_1 \text{ m } e_2$
- (22c) $e_2 \text{ di } e_3$
- (22d) $e_5 \text{ s } e_7$
- (22e) $e_8 > e_5$

Let us cluster the fluents as follows: for each fluent $a \in A$, fix a set $P = \{a\}$, and a set S whose members are these sets P .

$$(23) \quad S = \{\{e_1\}, \{e_2\}, \{e_3\}, \{e_4\}, \{e_5\}, \{e_6\}, \{e_7\}, \{e_8\}\}$$

Next, check for an Allen Relation between each pair of fluents a and a' . If a relation exists, add the fluent a' to the set P .

$$(24) \quad S' = \{\{e_1, e_4, e_2\}, \{e_2, e_3\}, \{e_3\}, \{e_4\}, \{e_5, e_7\}, \{e_6\}, \{e_7\}, \{e_8, e_5\}\}$$

Finally, for each pair of sets P, Q , if $|P \cap Q| > 0$, form $R = P \cup Q$, adding R to the set S'' and discarding P, Q . Add the remaining sets from S' to S'' .

$$(25) \quad S'' = \{\{e_1, e_4, e_2, e_3\}, \{e_5, e_7, e_8\}, \{e_6\}\}$$

We might form the underspecified event groups E_1 and E_2 to refer to these first two clusters, at which point we may freely treat these groups as normal bounded events, and perform asynchronous superposition on their event-strings, as well as with that of e_6 – reducing the number of inputs from 8 to 3.

Additionally, various weightings might be considered as a method of priority-ordering in the case of a large A , such as the number of component events in an underspecified event group, or the number of relations linking to a particular event.

8 Conclusion

We have explored in this work the possibility of using strings as models for event data, motivated by their nature as computational entities. The operation of asynchronous superposition was described for composing strings which represent finite, bounded events, as well as its limits in terms of massive growth in the number of outputs when the operation is repeated in sequence. The problem is addressed by constraining the strings which may be superposed, with the 13 unique Allen Relations forming the main part of these, as these can be found in annotated corpora such as TIMEBANK, using the ISO-TimeML standard.

Future work on this topic will further develop the constraints on asynchronous superposition and expand on the features of the current framework, while also examining the use of alternative models to approach the same issue, such as using finite state automata, or a hybrid string/FSA approach. We will additionally explore the potential of employing methods from distributed computing in order to tackle the combinatorial explosion that occurs in asynchronously superposing unconstrained bounded events.

Acknowledgements

This research is supported by Science Foundation Ireland (SFI) through the CNGL Programme (Grant 12/CE/I2267) in the ADAPT Centre (<https://www.adaptcentre.ie>) at Trinity College Dublin. The ADAPT Centre for Digital Content Technology is funded under the SFI Research Centres Programme (Grant 13/RC/2106) and is co-funded under the European Regional Development Fund.

References

- Allen, J. F. (1983). Maintaining Knowledge About Temporal Intervals. *Communications of the ACM* 26(11), 832–843.
- Fernando, T. (2015). The Semantics of Tense and Aspect: A Finite-State Perspective. In S. Lappin and C. Fox (Eds.), *The Handbook of Contemporary Semantic Theory*, pp. 203–236. John Wiley & Sons.
- Fernando, T. (2016). On Regular Languages Over Power Sets. *Journal of Language Modelling* 4(1), 29–56.
- Pustejovsky, J., P. Hanks, R. Sauri, A. See, R. Gaizauskas, A. Setzer, D. Radev, B. Sundheim, D. Day, L. Ferro, et al. (2003). The TIMEBANK Corpus. In *Corpus Linguistics*, Volume 2003, pp. 647–656. Lancaster, UK.
- Pustejovsky, J., K. Lee, H. Bunt, and L. Romary (2010). ISO-TimeML: An International Standard for Semantic Annotation. In *LREC*, Volume 10, pp. 394–397.
- Setzer, A., R. Gaizauskas, and M. Hepple (2005). The Role of Inference in the Temporal Annotation and Analysis of Text. *Language Resources and Evaluation* 39(2), 243–265.
- TimeML Working Group (2005). TimeML 1.2.1. A Formal Specification Language for Events and Temporal Expressions. http://www.timeml.org/publications/timeMLdocs/timeml_1.2.1.html#tlink. Accessed: 2017-07-12.